Scalable Machine Learning and Deep Learning - Review Questions 3

Deadline: November 22, 2020

1. **1 point.** Is it OK to initialize all the weights of a neural network to the same value as long as that value is selected randomly using He initialization? Is it okay to initialize the bias terms to 0?

---

2. **0.5 point.** In which cases would you want to use each of the following activation functions: ELU, leaky ReLU, ReLU, tanh, logistic, and softmax?

---

3. **0.5 point.** What is batch normalization and why does it work?

---

4. **1 point.** Does dropout slow down training? Does it slow down inference (i.e., making predictions on new instances)?

---

5. **1 point.** Consider a CNN composed of three convolutional layers, each with 3×3 filters, a stride of 2, and `SAME` padding. The lowest layer outputs 100 feature maps, the middle one outputs 200, and the top one outputs 400. The input images are RGB images of 200×300 pixels. What is the total number of parameters `w` in the CNN?

---

6. **1 point.** Consider a CNN with one convolutional layer, in which it has a 3×3 filter (as shown below) and a stride of 2. Please write the output of this layer for the given input image (the left image in the following figure)?