



# Machine Learning with Spark

Amir H. Payberah  
payberah@kth.se  
02/11/2018

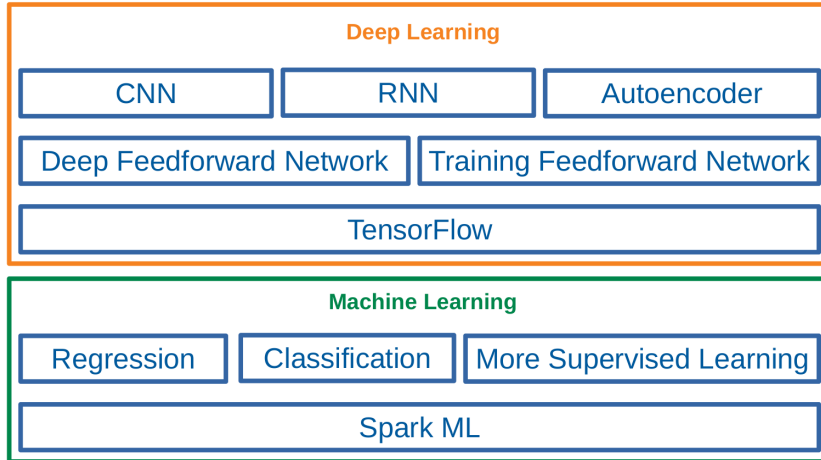




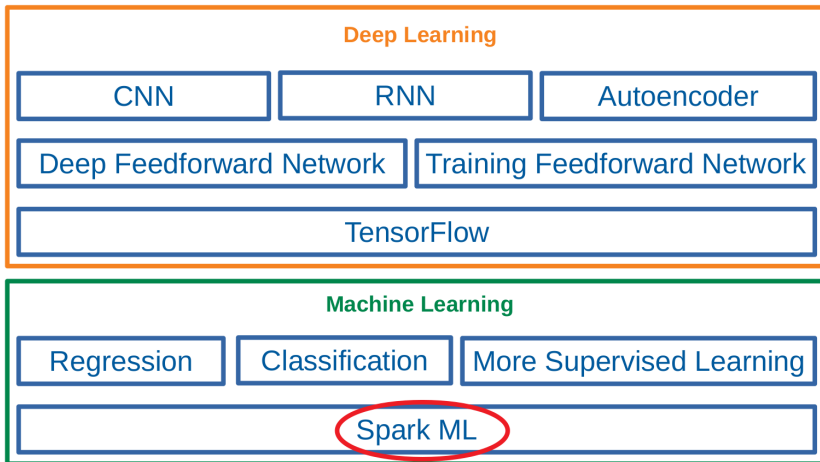
## The Course Web Page

<https://id2223kth.github.io>

# Where Are We?



## Where Are We?







# Problem

- ▶ Traditional platforms fail to show the expected performance.
- ▶ Need new systems to store and process large-scale data

# Scale Up vs. Scale Out

- ▶ Scale **up** or scale **vertically**
- ▶ Scale **out** or scale **horizontally**




# Spark




# Word Count

- ▶ Count the number of times each **distinct word** appears in the file
- ▶ If the file **fits in memory**: `words(doc.txt) | sort | uniq -c`



# Word Count

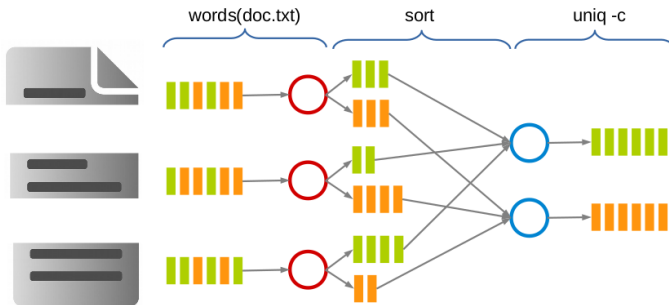
- ▶ Count the number of times each distinct word appears in the file
- ▶ If the file fits in memory: `words(doc.txt) | sort | uniq -c`



- ▶ If not?

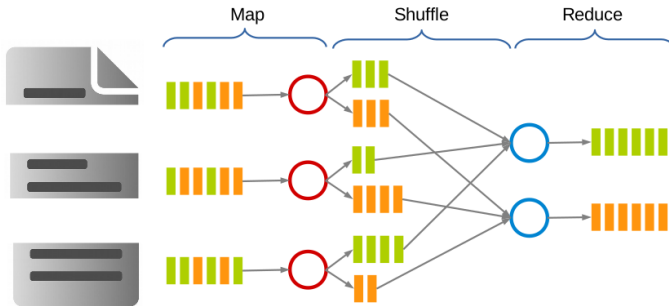
# Data-Parallel Processing (1/4)

- **Parallelize** the data and process.



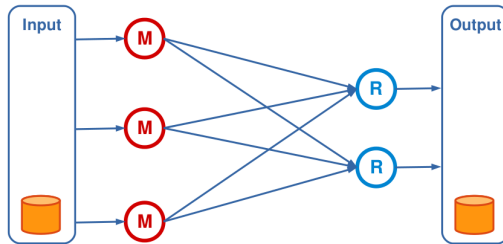
## Data-Parallel Processing (2/4)

### ► MapReduce




## Data-Parallel Processing (3/4)

- ▶ **Data flow** programming model.
- ▶ **Acyclic Directed Graph (DAG)** of **data flow** from stable storage to stable storage.




## Data-Parallel Processing (4/4)



# Spark Execution Model (1/3)


- ▶ Spark applications consist of
  - A driver process
  - A set of executor processes



[M. Zaharia et al., Spark: The Definitive Guide, O'Reilly Media, 2018]

## Spark Execution Model (2/3)

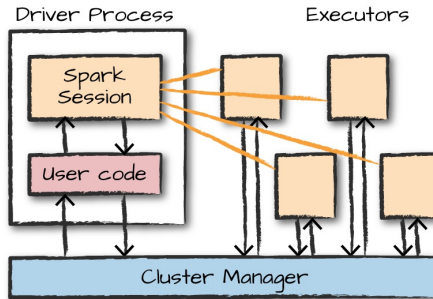
- ▶ The **driver process** is the **heart** of a **Spark application**
- ▶ Sits on a **node** in the cluster
- ▶ Runs the **main()** function





## Spark Execution Model (3/3)

- **Executors** execute codes assigned to them by the **driver**.



## Resilient Distributed Datasets (RDD) (1/2)

- ▶ A distributed memory abstraction.
- ▶ Immutable collections of objects spread across a cluster.
  - Like a `LinkedList <MyObjects>`



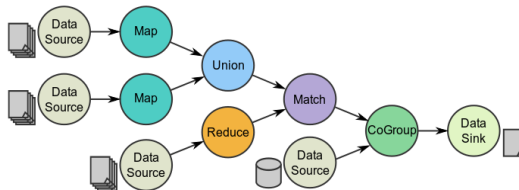
## Resilient Distributed Datasets (RDD) (2/2)

- ▶ An **RDD** is divided into a number of **partitions**, which are **atomic** pieces of information.
- ▶ Partitions of an RDD can be stored on different **nodes** of a cluster.



# Spark Programming Model

- ▶ **Job** description based on **directed acyclic graphs (DAG)**.
- ▶ There are two types of RDD operators: **transformations** and **actions**.






# RDD Operations

- ▶ **Transformations:** *lazy* operators that create *new* RDDs.
- ▶ **Actions:** launch a *computation* and return a *value* to the program or write data to the external storage.

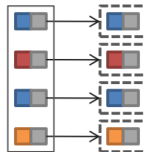
## RDD Transformations (1/3)

- **Map:** all pairs are **independently** processed.



## RDD Transformations (1/3)

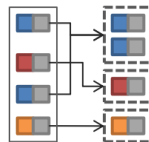
- **Map:** all pairs are **independently** processed.



```
// passing each element through a function.  
val nums = sc.parallelize(Array(1, 2, 3))  
val squares = nums.map(x => x * x) // {1, 4, 9}  
  
// selecting those elements that func returns true.  
val even = squares.filter(x => x % 2 == 0) // {4}  
  
// mapping each element to zero or more others.  
nums.flatMap(x => Range(0, x)) // {0, 0, 1, 0, 1, 2}
```

## RDD Transformations (2/3)


- ▶ **Group by:** pairs with **identical key** are grouped.
- ▶ Groups are independently processed.





## RDD Transformations (2/3)

- ▶ **Group by:** pairs with **identical key** are grouped.
- ▶ Groups are independently processed.



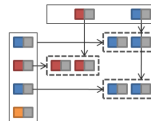
```
val pets = sc.parallelize(Seq(("kth", 1), ("rise", 1), ("kth", 2)))

pets.groupByKey()
// {(kth, (1, 2)), (rise, (1))}

pets.reduceByKey((x, y) => x + y)
// {(kth, 3), (rise, 1)}
```

## RDD Transformations (3/3)

- ▶ **Join**: performs an **equi-join** on the key.
- ▶ Join candidates are independently processed.



- 

```
rdd1.join(rdd2)
// ("a", (1, "g1"))
// ("a", (3, "g1"))
// ("b", (2, "g2"))
```



## Basic RDD Actions (1/2)

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))  
nums.collect() // Array(1, 2, 3)
```



## Basic RDD Actions (1/2)

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))  
nums.collect() // Array(1, 2, 3)
```

- ▶ Return an array with the first n elements of the RDD.

```
nums.take(2) // Array(1, 2)
```



## Basic RDD Actions (1/2)

- ▶ Return all the elements of the RDD as an array.

```
val nums = sc.parallelize(Array(1, 2, 3))  
nums.collect() // Array(1, 2, 3)
```

- ▶ Return an array with the first n elements of the RDD.

```
nums.take(2) // Array(1, 2)
```

- ▶ Return the number of elements in the RDD.

```
nums.count() // 3
```



## Basic RDD Actions (2/2)

- ▶ Aggregate the elements of the RDD using the given function.

```
nums.reduce((x, y) => x + y)  
// or  
nums.reduce(_ + _) // 6
```



## Basic RDD Actions (2/2)

- ▶ Aggregate the elements of the RDD using the given function.

```
nums.reduce((x, y) => x + y)  
// or  
nums.reduce(_ + _) // 6
```

- ▶ Write the elements of the RDD as a text file.

```
nums.saveAsTextFile("hdfs://file.txt")
```





## Creating RDDs

- ▶ Turn a collection into an RDD.

```
val a = sc.parallelize(Array(1, 2, 3))
```



## Creating RDDs

- ▶ Turn a collection into an RDD.

```
val a = sc.parallelize(Array(1, 2, 3))
```

- ▶ Load text file from local FS, HDFS, or S3.

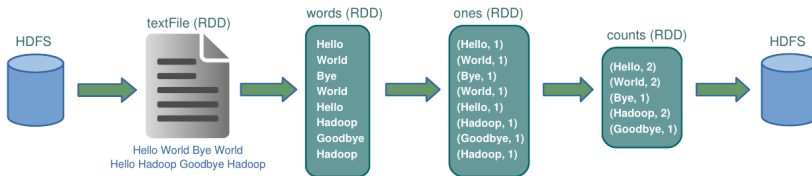
```
val a = sc.textFile("file.txt")  
val b = sc.textFile("directory/*.txt")  
val c = sc.textFile("hdfs://namenode:9000/path/file")
```

# Example

```
val textFile = sc.textFile("hdfs://...")

val words = textFile.flatMap(line => line.split(" "))
val ones = words.map(word => (word, 1))
val counts = ones.reduceByKey(_ + _)

counts.saveAsTextFile("hdfs://...")
```




## Example

```
val textFile = sc.textFile("hdfs://...")

val words = textFile.flatMap(line => line.split(" "))
val ones = words.map(word => (word, 1))
val counts = ones.reduceByKey(_ + _)

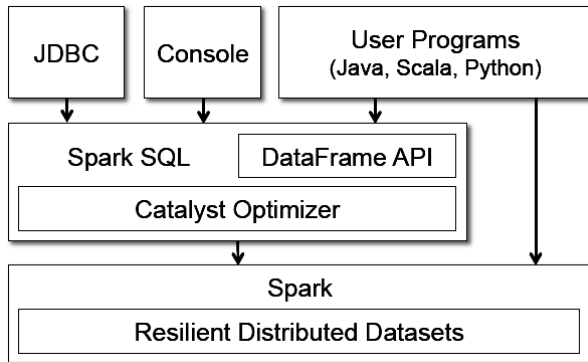
counts.saveAsTextFile("hdfs://...")
```



► **Lineage:** transformations used to build an RDD.

# Spark SQL

## Spark and Spark SQL



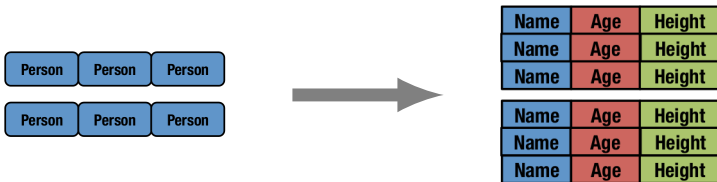


# DataFrame

- ▶ A **DataFrame** is a distributed collection of rows with a homogeneous schema.
- ▶ It is equivalent to a **table** in a relational database.
- ▶ It can also be manipulated in similar ways to **RDDs**.

## Adding Schema to RDDs

- ▶ **Spark + RDD**: **functional** transformations on partitioned collections of **opaque objects**.
- ▶ **SQL + DataFrame**: **declarative** transformations on partitioned collections of **tuples**.







## Creating a DataFrame - From an RDD

- You can use `toDF` to convert an RDD to DataFrame.

```
val tupleRDD = sc.parallelize(Array(("seif", 65, 0), ("amir", 40, 1)))  
val tupleDF = tupleRDD.toDF("name", "age", "id")
```



## Creating a DataFrame - From an RDD

- ▶ You can use `toDF` to convert an RDD to DataFrame.

```
val tupleRDD = sc.parallelize(Array(("seif", 65, 0), ("amir", 40, 1)))  
val tupleDF = tupleRDD.toDF("name", "age", "id")
```

- ▶ If RDD contains `case` class instances, Spark infers the attributes from it.

```
case class Person(name: String, age: Int, id: Int)  
  
val peopleRDD = sc.parallelize(Array(Person("seif", 65, 0), Person("amir", 40, 1)))  
val peopleDF = peopleRDD.toDF
```



# Creating a DataFrame - From Data Source

- ▶ Data sources supported by Spark.
  - CSV, JSON, Parquet, ORC, JDBC/ODBC connections, Plain-text files
  - Cassandra, HBase, MongoDB, AWS Redshift, XML, etc.

```
val peopleJson = spark.read.format("json").load("people.json")

val peopleCsv = spark.read.format("csv")
  .option("sep", ";")
  .option("inferSchema", "true")
  .option("header", "true")
  .load("people.csv")
```

- Different ways to refer to a column.

```
val people = spark.read.format("json").load("people.json")  
  
people.col("name")  
  
col("name")  
  
column("name")  
  
'name  
  
$"name"  
  
expr("name")
```



## Row

- ▶ A **row** is a **record of data**.
- ▶ They are of type **Row**.
- ▶ Rows do **not have schemas**.
- ▶ To access data in rows, you need to specify the **position** that you would like.

```
import org.apache.spark.sql.Row

val myRow = Row("Seif", 65, 0)

myRow(0) // type Any
myRow(0).asInstanceOf[String] // String
myRow.getString(0) // String
myRow.getInt(1) // Int
```



## DataFrame Transformations (1/6)

- ▶ `select` allows to do the **DataFrame equivalent** of **SQL queries** on a table of data.

```
people.select("name", "age", "id").show(2)
people.select(col("name"), expr("age + 3")).show()
people.select(expr("name AS username")).show(2)
```



## DataFrame Transformations (1/6)

- ▶ `select` allows to do the **DataFrame equivalent** of **SQL queries** on a table of data.

```
people.select("name", "age", "id").show(2)
people.select(col("name"), expr("age + 3")).show()
people.select(expr("name AS username")).show(2)
```

- ▶ `filter` and `where` both **filter** rows.

```
people.filter(col("age") < 20).show()
people.where("age < 20").show()
```



## DataFrame Transformations (1/6)

- ▶ `select` allows to do the **DataFrame equivalent** of **SQL queries** on a table of data.

```
people.select("name", "age", "id").show(2)
people.select(col("name"), expr("age + 3")).show()
people.select(expr("name AS username")).show(2)
```

- ▶ `filter` and `where` both **filter** rows.

```
people.filter(col("age") < 20).show()
people.where("age < 20").show()
```

- ▶ `distinct` can be used to extract unique rows.

```
people.select("name").distinct().count()
```





## DataFrame Transformations (2/6)

- ▶ `withColumn` adds a new column to a DataFrame.

```
people.withColumn("teenager", expr("age < 20")).show()
```



## DataFrame Transformations (2/6)

- ▶ `withColumn` adds a new column to a DataFrame.

```
people.withColumn("teenager", expr("age < 20")).show()
```

- ▶ `withColumnRenamed` renames a column.

```
people.withColumnRenamed("name", "username").columns
```



## DataFrame Transformations (2/6)

- ▶ `withColumn` adds a new column to a DataFrame.

```
people.withColumn("teenager", expr("age < 20")).show()
```

- ▶ `withColumnRenamed` renames a column.

```
people.withColumnRenamed("name", "username").columns
```

- ▶ `drop` removes a column.

```
people.drop("name").columns
```



## DataFrame Transformations (3/6)

- ▶ `count` returns the total number of values.

```
people.select(count("age")).show()
```



## DataFrame Transformations (3/6)

- ▶ `count` returns the total number of values.

```
people.select(count("age")).show()
```

- ▶ `countDistinct` returns the number of unique groups.

```
people.select(countDistinct("name")).show()
```



## DataFrame Transformations (3/6)

- ▶ `count` returns the **total number of values**.

```
people.select(count("age")).show()
```

- ▶ `countDistinct` returns the **number of unique groups**.

```
people.select(countDistinct("name")).show()
```

- ▶ `first` and `last` return the **first and last value** of a DataFrame.

```
people.select(first("name"), last("age")).show()
```



## DataFrame Transformations (4/6)

- ▶ `min` and `max` extract the **minimum and maximum values** from a DataFrame.

```
people.select(min("name"), max("age"), max("id")).show()
```



## DataFrame Transformations (4/6)

- ▶ `min` and `max` extract the **minimum and maximum values** from a DataFrame.

```
people.select(min("name"), max("age"), max("id")).show()
```

- ▶ `sum` **adds all the values** in a column.

```
people.select(sum("age")).show()
```





## DataFrame Transformations (4/6)

- ▶ `min` and `max` extract the **minimum and maximum values** from a DataFrame.

```
people.select(min("name"), max("age"), max("id")).show()
```

- ▶ `sum` adds all the values in a column.

```
people.select(sum("age")).show()
```

- ▶ `avg` calculates the **average**.

```
people.select(avg("age")).show()
```



## DataFrame Transformations (5/6)

- ▶ `groupBy` and `agg` together perform aggregations on `groups`.

```
people.groupBy("name").agg(count("age")).show()
```

## DataFrame Transformations (5/6)

- ▶ `groupBy` and `agg` together perform aggregations on `groups`.

```
people.groupBy("name").agg(count("age")).show()
```

- ▶ `join` performs the join operation between `two tables`.

```
val t1 = Seq((0, "a", 0), (1, "b", 1), (2, "c", 1)).toDF("num", "name", "id")
val t2 = Seq((0, "x"), (1, "y"), (2, "z")).toDF("id", "group")

val joinExpression = t1.col("id") === t2.col("id")
var joinType = "inner"

t1.join(t2, joinExpression, joinType).show()
```

## DataFrame Transformations (6/6)

- ▶ You can use `udf` to define new **column-based functions**.

```
import org.apache.spark.sql.functions.udf

val df = Seq((0, "hello"), (1, "world")).toDF("id", "text")

val upper: String => String = _.toUpperCase

val upperUDF = udf(upper)

df.withColumn("upper", upperUDF(col("text"))).show
```



## DataFrame Actions

- ▶ Like RDDs, DataFrames also have their own set of actions.
- ▶ `collect`: returns an `array` that contains all the `rows` in this DataFrame.
- ▶ `count`: returns the `number of rows` in this DataFrame.
- ▶ `first` and `head`: returns the `first row` of the DataFrame.
- ▶ `show`: displays the `top 20 rows` of the DataFrame in a tabular form.
- ▶ `take`: returns the `first n rows` of the DataFrame.

# Machine Learning



# Machine Learning with Spark

- ▶ Spark provides support for **statistics** and **machine learning**.
  - Supervised learning
  - Unsupervised engines
  - Deep learning



# Supervised Learning

- ▶ Using **labeled historical data** and **training a model** to **predict** the values of those labels **based on various features** of the data points.
- ▶ **Classification** (**categorical** values)
  - E.g., predicting disease, classifying images, ...
- ▶ **Regression** (**continuous** values)
  - E.g., predicting sales, predicting height, ...



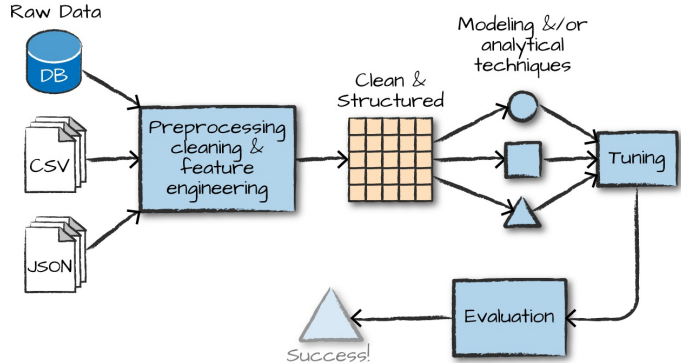


# Unsupervised Learning

- ▶ No label to predict.
- ▶ Trying to find patterns or discover the underlying structure in a given set of data.
  - Clustering, anomaly detection, ...

# The Advanced Analytic Process

- ▶ Data collection
- ▶ Data cleaning
- ▶ Feature engineering
- ▶ Training models
- ▶ Model tuning and evaluation





## What is MLlib? (1/2)

- ▶ **MLlib** is a package built on **Spark**.
- ▶ It provides **interfaces** for:
  - **Gathering** and **cleaning** data
  - **Feature engineering** and feature selection
  - **Training** and **tuning** large-scale **supervised and unsupervised** machine learning models
  - Using those models in **production**



## What is MLlib? (2/2)

- ▶ MLlib consists of **two packages**.



## What is MLlib? (2/2)

- ▶ MLlib consists of `two packages`.
- ▶ `org.apache.spark.mllib`
  - Uses `RDDs`
  - It is in `maintenance mode` (only receives `bug fixes`, not new features)



## What is MLlib? (2/2)

- ▶ MLlib consists of `two packages`.
- ▶ `org.apache.spark.mllib`
  - Uses `RDDs`
  - It is in `maintenance mode` (only receives `bug fixes`, not new features)
- ▶ `org.apache.spark.ml`
  - Uses `DataFrames`
  - Offers a `high-level` interface for building `machine learning pipelines`



## When Should We Use MLlib?

- ▶ Many tools for performing machine learning on a single machine, e.g., scikit-learn and TensorFlow.



## When Should We Use MLlib?

- ▶ Many tools for performing machine learning on a single machine, e.g., scikit-learn and TensorFlow.
- ▶ These single-machine tools are usually complementary to MLlib.



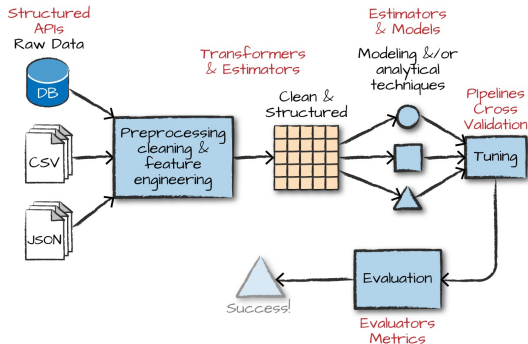


## When Should We Use MLlib?

- ▶ Many tools for performing machine learning on a single machine, e.g., scikit-learn and TensorFlow.
- ▶ These single-machine tools are usually complementary to MLlib.
- ▶ Take advantage of Spark, when you hit scalability issues.
  - Use Spark for preprocessing and feature generation, before giving data to single-machine learning libraries.
  - Use Spark, when input data or model size become too difficult to put on one machine.

# High-Level MLlib Concepts

- ▶ ML Pipelines ([spark.ml](https://spark.apache.org/docs/latest/ml-pipeline.html)) provide a **uniform set of high-level APIs** built on top of **DataFrames** to create **machine learning pipelines**.
- ▶ Main pipeline components: **transformers** and **estimators**

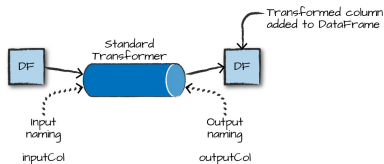



# Transformers

- ▶ **Transformers** take a **DataFrame** as input and produce a new **DataFrame** as output.
- ▶ The class **Transformer** implements a method **transform()** that converts **one DataFrame into another**.

```
// transformer: DataFrame => DataFrame
```

```
transform(dataset: DataFrame): DataFrame
```

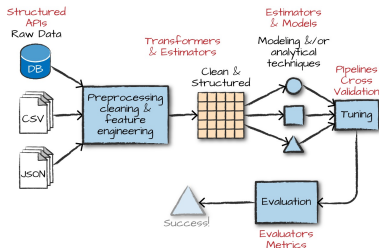


# Estimators

- **Estimator** is an abstraction of a learning algorithm that fits a **model** on a **dataset**.
- The class **Estimator** implements a method **fit()**, which accepts a **DataFrame** and produces a **Model (Transformer)**.

```
// estimator: DataFrame => Model
```

```
fit(dataset: DataFrame): M
```





# Pipeline

- ▶ **Pipeline** is a sequence of algorithms to process and learn from data.



# Pipeline


- ▶ **Pipeline** is a **sequence of algorithms** to **process and learn** from data.
- ▶ E.g., a **text document processing workflow** might include several **stages**:
  - **Split** each document's **text into words**.
  - **Convert** each document's **words into a numerical feature vector**.
  - Learn a **prediction model** using the feature vectors and labels.

# Pipeline

- ▶ **Pipeline** is a **sequence of algorithms** to **process and learn** from data.
- ▶ E.g., a **text document processing workflow** might include several **stages**:
  - **Split** each document's **text into words**.
  - **Convert** each document's **words into a numerical feature vector**.
  - Learn a **prediction model** using the feature vectors and labels.
- ▶ MLlib **represents such a workflow** as a **Pipeline**, which is a **sequence of stages**.
  - Each stage is either a **Transformer** or an **Estimator**.

# How Does Pipeline Work? (1/3)

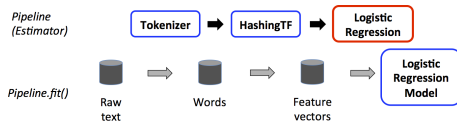
- ▶ Stages of a pipeline run in order.
- ▶ The input `DataFrame` is transformed as it passes through each stage.
  - Calls `transform()` for `Transformer` stages
  - Calls `fit()` for `Estimator` stages





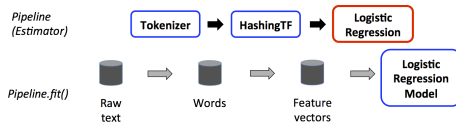
# How Does Pipeline Work? (1/3)

- ▶ Stages of a pipeline run in order.
- ▶ The input `DataFrame` is transformed as it passes through each stage.
  - Calls `transform()` for `Transformer` stages
  - Calls `fit()` for `Estimator` stages
- ▶ E.g., a Pipeline with three stages: `Tokenizer` and `HashingTF` are `Transformers`, and `LogisticRegression` is an `Estimator`.




## How Does Pipeline Work? (2/3)

- ▶ `Pipeline.fit()`: is called on the **original DataFrame**
  - DataFrame with **raw text documents and labels**




## How Does Pipeline Work? (2/3)

- ▶ `Pipeline.fit()`: is called on the **original DataFrame**
  - DataFrame with **raw text documents and labels**
- ▶ `Tokenizer.transform()`: **splits the raw text** documents into words
  - Adds a **new column with words** to the DataFrame




## How Does Pipeline Work? (2/3)

- ▶ `Pipeline.fit()`: is called on the **original DataFrame**
  - DataFrame with **raw text documents and labels**
- ▶ `Tokenizer.transform()`: **splits the raw text** documents into words
  - Adds a **new column with words** to the DataFrame
- ▶ `HashingTF.transform()`: **converts the words** column into **feature vectors**
  - Adds **new column with those vectors** to the DataFrame




## How Does Pipeline Work? (2/3)

- ▶ `Pipeline.fit()`: is called on the **original DataFrame**
  - DataFrame with **raw text documents and labels**
- ▶ `Tokenizer.transform()`: **splits the raw text** documents into words
  - Adds a **new column with words** to the DataFrame
- ▶ `HashingTF.transform()`: **converts the words** column into **feature vectors**
  - Adds **new column with those vectors** to the DataFrame
- ▶ `LogisticRegression.fit()`: produces a **model** (`LogisticRegressionModel`).




# How Does Pipeline Work? (3/3)

- ▶ A **Pipeline** is an **Estimator** (`DataFrame`  $\Rightarrow$  `Model`).



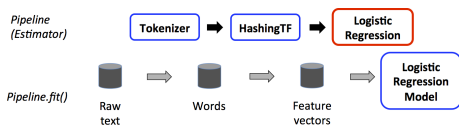
Pipeline



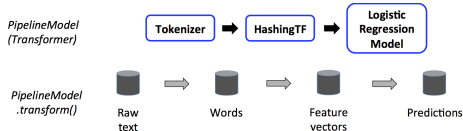
PipelineModel

## How Does Pipeline Work? (3/3)

- ▶ A **Pipeline** is an **Estimator** (`DataFrame`  $\Rightarrow$  `Model`).
- ▶ After a **Pipeline**'s `fit()` runs, it produces a **PipelineModel**.




Pipeline




PipelineModel

## How Does Pipeline Work? (3/3)

- ▶ A **Pipeline** is an **Estimator** (`DataFrame`  $\Rightarrow$  `Model`).
- ▶ After a **Pipeline**'s `fit()` runs, it **produces** a **PipelineModel**.
- ▶ **PipelineModel** is a **Transformer** (`DataFrame`  $\Rightarrow$  `DataFrame`).



Pipeline

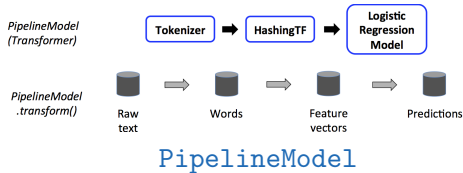
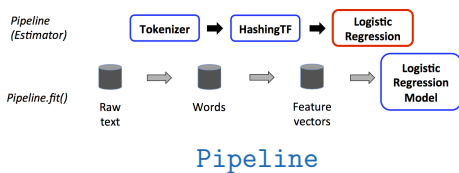


PipelineModel





## How Does Pipeline Work? (3/3)

- ▶ A **Pipeline** is an **Estimator** (`DataFrame = [fit] => Model`).
- ▶ After a **Pipeline**'s `fit()` runs, it **produces** a **PipelineModel**.
- ▶ **PipelineModel** is a **Transformer** (`DataFrame = [transform] => DataFrame`).
- ▶ The **PipelineModel** is used at **test time**.



# How Does Pipeline Work? (3/3)

- ▶ A **Pipeline** is an **Estimator** (`DataFrame = [fit] => Model`).
- ▶ After a **Pipeline**'s `fit()` runs, it produces a **PipelineModel**.
- ▶ **PipelineModel** is a **Transformer** (`DataFrame = [transform] => DataFrame`).
- ▶ The **PipelineModel** is used at **test time**.
- ▶ The **PipelineModel** has the **same number of stages** as the original **Pipeline**.
  - All **Estimators** in the original **Pipeline** have become **Transformers**.





# Parameters

- ▶ MLlib `Estimators` and `Transformers` use a `uniform API` for `specifying parameters`.



# Parameters

- ▶ MLlib **Estimators** and **Transformers** use a **uniform API** for **specifying parameters**.
- ▶ **Param**: a **named parameter**
- ▶ **ParamMap**: a set of **(parameter, value)** pairs



# Parameters

- ▶ MLlib `Estimators` and `Transformers` use a `uniform API` for `specifying parameters`.
- ▶ `Param`: a `named parameter`
- ▶ `ParamMap`: a set of `(parameter, value)` pairs
- ▶ Two ways to `pass parameters` to an algorithm:
  1. Set parameters for an instance, e.g., `ls.setMaxIter(10)`
  2. Pass a `ParamMap` to `fit()` or `transform()`.



## Example - Input DataFrame (1/2)

- Make a DataFrame of the type `Article`.

```
import spark.implicits._

case class Article(id: Long, topic: String, text: String)

val articles = Seq(Article(0, "sci.math", "Hello, Math!"),
                   Article(1, "alt.religion", "Hello, Religion!"),
                   Article(2, "sci.physics", "Hello, Physics!"),
                   Article(3, "sci.math", "Hello, Math Revised!"),
                   Article(4, "sci.math", "Better Math"),
                   Article(5, "alt.religion", "TGIF")).toDF

articles.show
```



## Example - Input DataFrame (2/2)

- ▶ Add a new column `label` to the DataFrame.
- ▶ `udf` is a feature of Spark SQL to define new Column-based functions.

```
val topic2Label: Boolean => Double = x => if (x) 1 else 0  
  
val toLabel = udf(topic2Label)  
  
val labelled = articles.withColumn("label", toLabel($"topic".like("sci%"))).cache
```



## Example - Transformers (1/2)

- Break each sentence into individual terms (words).

```
import org.apache.spark.ml.feature.Tokenizer
import org.apache.spark.ml.feature.RegexTokenizer

val tokenizer = new RegexTokenizer().setInputCol("text").setOutputCol("words")

val tokenized = tokenizer.transform(labelled)

tokenized.show
```



## Example - Transformers (2/2)

- ▶ Takes a set of words and converts them into **fixed-length feature vector**.
  - 5000 in our example
- ▶ Uses a **hash function** to map each word into an **index** in the feature vector.
- ▶ Then computes the **term frequencies** based on the mapped indices.

```
import org.apache.spark.ml.feature.HashingTF

val hashingTF = new HashingTF().setInputCol(tokenizer.getOutputCol)
                                .setOutputCol("features")
                                .setNumFeatures(5000)

val hashed = hashingTF.transform(tokenized)

hashed.show
```



## Example - Estimator

```
val Array(trainDF, testDF) = hashed.randomSplit(Array(0.8, 0.2))
```

```
trainDF.show
```

```
testDF.show
```

```
import org.apache.spark.ml.classification.LogisticRegression
```

```
val lr = new LogisticRegression().setMaxIter(20).setRegParam(0.01)
```

```
val model = lr.fit(trainDF)
```

```
val pred = model.transform(testDF).select("topic", "label", "prediction")
```

```
pred.show
```

## Example - Pipeline

```
val Array(trainDF2, testDF2) = labelled.randomSplit(Array(0.8, 0.2))
```

```
trainDF2.show
```

```
testDF2.show
```

```
import org.apache.spark.ml.{Pipeline, PipelineModel}
```

```
val pipeline = new Pipeline().setStages(Array(tokenizer, hashingTF, lr))
```

```
val model2 = pipeline.fit(trainDF2)
```

```
val pred = model2.transform(testDF2).select("topic", "label", "prediction")
```

```
pred.show
```



## Example - ParamMap

```
// set parameters using setter methods.
```

```
val lr = new LogisticRegression()
```

```
lr.setMaxIter(10).setRegParam(0.01)
```

```
// specify parameters using a ParamMap
```

```
val lr = new LogisticRegression()
```

```
val paramMap = ParamMap(lr.maxIter -> 20)
```

```
.put(lr.maxIter, 30) // specify one Param
```

```
.put(lr.regParam -> 0.1, lr.threshold -> 0.55) // specify multiple Params
```

```
val model = lr.fit(training, paramMap)
```



## Low-Level Data Types - Local Vector

- ▶ Stored on a **single** machine
- ▶ **Dense** and **sparse**
  - **Dense** (1.0, 0.0, 3.0): [1.0, 0.0, 3.0]
  - **Sparse** (1.0, 0.0, 3.0): (3, [0, 2], [1.0, 3.0])

```
import org.apache.spark.mllib.linalg.{Vector, Vectors}

val dv: Vector = Vectors.dense(1.0, 0.0, 3.0)

val sv1: Vector = Vectors.sparse(3, Array(0, 2), Array(1.0, 3.0))
val sv2: Vector = Vectors.sparse(3, Seq((0, 1.0), (2, 3.0)))
```


# Preprocessing and Feature Engineering




# Formatting Models

- ▶ In most of **classification and regression** algorithms, we want to get the **data**.
  - A **column** to represent the **label** (**Double**).
  - A **column** to represent the **features** (**Vector**)

# Transformers and Estimators



Transformer



Estimator





# Transformer Properties

- ▶ All transformers require you to specify the **input and output columns**.
- ▶ We can set these with **setInputCol** and **setOutputCol**.

```
val tokenizer = new RegexTokenizer().setInputCol("text").setOutputCol("words")
```



# Vector Assembler

- ▶ Concatenate all your features into one vector.

```
import org.apache.spark.ml.feature.VectorAssembler

case class Nums(val1: Long, val2: Long, val3: Long)

val numsDF = Seq(Nums(1, 2, 3), Nums(4, 5, 6), Nums(7, 8, 9)).toDF

val va = new VectorAssembler().setInputCols(Array("val1", "val2", "val3"))
                                .setOutputCol("features")

va.transform(numsDF).show()
```



# MLlib Transformers

- ▶ Continuous features
- ▶ Categorical features
- ▶ Text data



# MLlib Transformers

- ▶ Continuous features
- ▶ Categorical features
- ▶ Text data

## Continuous Features - Bucketing

- Convert continuous features into categorical features.

```
import org.apache.spark.ml.feature.Bucketizer

val contDF = spark.range(20).selectExpr("cast(id as double)")
val bucketBorders = Array(-1.0, 5.0, 10.0, 15.0, 20.0)

val bucketer = new Bucketizer().setSplits(bucketBorders).setInputCol("id")

bucketer.transform(contDF).show()
```

# Continuous Features - Scaling and Normalization

- To **scale** and **normalize** continuous data.

```
import org.apache.spark.ml.feature.VectorAssembler

case class Nums(val1: Long, val2: Long, val3: Long)
val numsDF = Seq(Nums(1, 2, 3), Nums(4, 5, 6), Nums(7, 8, 9)).toDF
val va = new VectorAssembler().setInputCols(Array("val1", "val2", "val3"))
                                   .setOutputCol("features")
val nums = va.transform(numsDF)
```

```
import org.apache.spark.ml.feature.StandardScaler

val scaler = new StandardScaler().setInputCol("features").setOutputCol("scaled")
scaler.fit(nums).transform(nums).show()
```

## Continuous Features - Maximum Absolute Scaler

- Scales the data by dividing each feature by the maximum absolute value in this feature (column).

```
import org.apache.spark.ml.feature.VectorAssembler

case class Nums(val1: Long, val2: Long, val3: Long)
val numsDF = Seq(Nums(1, 2, 3), Nums(4, 5, 6), Nums(7, 8, 9)).toDF
val va = new VectorAssembler().setInputCols(Array("val1", "val2", "val3"))
                                   .setOutputCol("features")
val nums = va.transform(numsDF)
```

```
import org.apache.spark.ml.feature.MaxAbsScaler

val maScaler = new MaxAbsScaler().setInputCol("features").setOutputCol("mas")
maScaler.fit(nums).transform(nums).show()
```



# MLlib Transformers

- ▶ Continuous features
- ▶ Categorical features
- ▶ Text data





## Categorical Features - String Indexer

- Maps **strings** to different **numerical IDs**.

```
val simpleDF = spark.read.json("simple-ml.json")
```

```
import org.apache.spark.ml.feature.StringIndexer

val lblIndxr = new StringIndexer().setInputCol("lab").setOutputCol("labelInd")
val idxRes = lblIndxr.fit(simpleDF).transform(simpleDF)

idxRes.show()
```

## Categorical Features - Converting Indexed Values Back to Text

- Maps back to the original values.

```
import org.apache.spark.ml.feature.IndexToString  
  
val labelReverse = new IndexToString().setInputCol("labelInd").setOutputCol("original")  
  
labelReverse.transform(idxRes).show()
```

# Categorical Features - One-Hot Encoding

- Converts each **distinct value** to a **boolean flag** as a component in a **vector**.

```
val simpleDF = spark.read.json("simple-ml.json")
```

```
import org.apache.spark.ml.feature.OneHotEncoder
```

```
val lblIndxr = new StringIndexer().setInputCol("color").setOutputCol("colorInd")
val colorLab = lblIndxr.fit(simpleDF).transform(simpleDF.select("color"))
val ohe = new OneHotEncoder().setInputCol("colorInd").setOutputCol("one-hot")
ohe.transform(colorLab).show()
```

```
// Since there are three values, the vector is of length 2 and the mapping is as follows:
// 0  -> 10, (2, [0], [1.0])
// 1  -> 01, (2, [1], [1.0])
// 2  -> 00, (2, [], [])
// (2, [0], [1.0]) means a vector of length 2 with 1.0 at position 0 and 0 elsewhere.
```



# MLlib Transformers

- ▶ Continuous features
- ▶ Categorical features
- ▶ Text data

# Text Data - Tokenizing Text

- Converting free-form text into a list of tokens or individual words.

```
val sales = spark.read.format("csv").option("header", "true").load("sales.csv")  
  .where("Description IS NOT NULL")  
  
sales.show(false)
```

```
import org.apache.spark.ml.feature.Tokenizer  
  
val tkn = new Tokenizer().setInputCol("Description").setOutputCol("DescOut")  
val tokenized = tkn.transform(sales.select("Description"))  
tokenized.show(false)
```

## Text Data - Removing Common Words

- Filters stop words, such as "the", "and", and "but".

```
import org.apache.spark.ml.feature.StopWordsRemover

val df = spark.createDataFrame(Seq((0, Seq("I", "saw", "the", "red", "balloon")),
  (1, Seq("Mary", "had", "a", "little", "lamb")))).toDF("id", "raw")

val englishStopWords = StopWordsRemover.loadDefaultStopWords("english")

val stops = new StopWordsRemover().setStopWords(englishStopWords)
  .setInputCol("raw").setOutputCol("WithoutStops")

stops.transform(df).show(false)
```

# Text Data - Converting Words into Numerical Representations

- ▶ Counts instances of words in word features.
- ▶ Treats every row as a document, every word as a term, and the total collection of all terms as the vocabulary.

```
import org.apache.spark.ml.feature.CountVectorizer

val df = spark.createDataFrame(Seq((0, Array("a", "b", "c")),
    (1, Array("a", "b", "b", "c", "a")))).toDF("id", "words")

val cvModel = new CountVectorizer().setInputCol("words").setOutputCol("features")
    .setVocabSize(3).setMinDF(2)

val fittedCV = cvModel.fit(df)

fittedCV.transform(df).show(false)
```



## Text Data - TF-IDF (1/2)

- ▶ **TF-IDF**: term frequency - inverse document frequency
- ▶ Measures how often a word occurs in each document, weighted according to how many documents that word occurs in.
- ▶ Words that occur in a few documents are given more weight than words that occur in many documents.



## Text Data - TF-IDF (2/2)

```
import org.apache.spark.ml.feature.{HashingTF, IDF, Tokenizer}

val sentenceData = spark.createDataFrame(Seq((0.0, "Hi I heard about Spark"),
      (0.0, "I wish Java could use case classes"),
      (1.0, "Logistic regression models are neat")))
      .toDF("label", "sentence")

val tokenizer = new Tokenizer().setInputCol("sentence").setOutputCol("words")
val wordsData = tokenizer.transform(sentenceData)

val hashingTF = new HashingTF().setInputCol("words").setOutputCol("rawFeatures")
      .setNumFeatures(20)
val featurizedData = hashingTF.transform(wordsData)

val idf = new IDF().setInputCol("rawFeatures").setOutputCol("features")
val idfModel = idf.fit(featurizedData)

val rescaledData = idfModel.transform(featurizedData)
rescaledData.select("label", "features").show(false)
```

# Summary



# Summary

- ▶ Spark: RDD
- ▶ Spark SQL: DataFrame
- ▶ MLlib
  - Transformers and Estimators
  - Pipeline
  - Feature engineering



## References

- ▶ Matei Zaharia et al., Spark - The Definitive Guide, (Ch. 24 and 25)

Questions?