



# TensorFlow

Amir H. Payberah  
payberah@kth.se  
23/11/2018

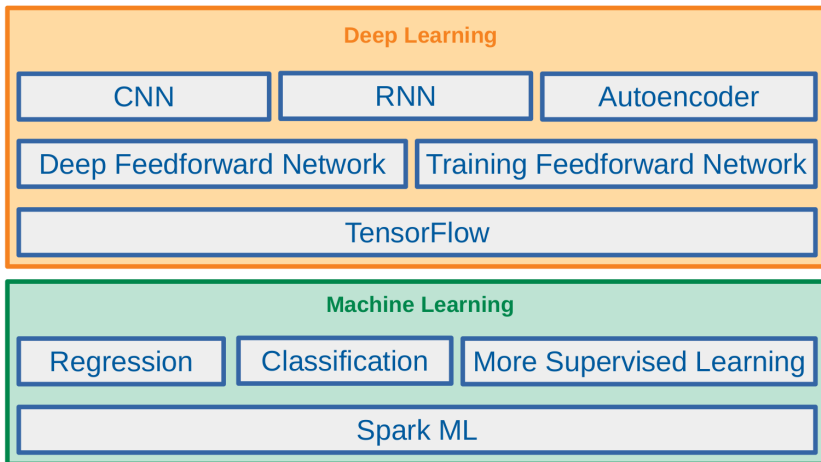




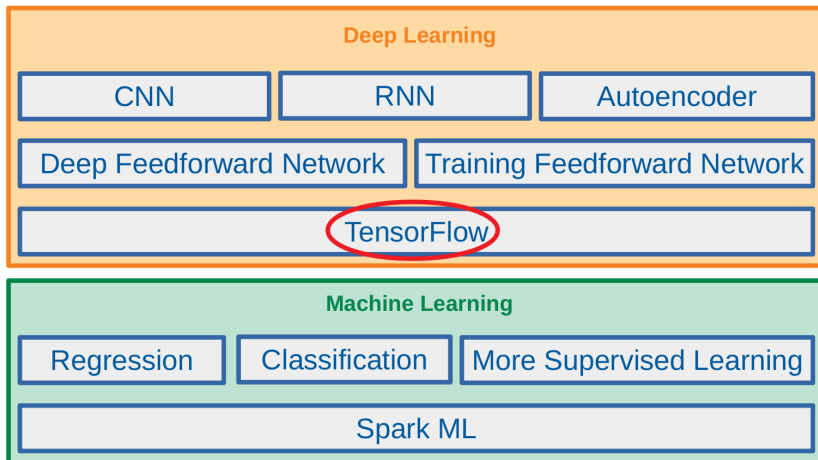
## The Course Web Page

<https://id2223kth.github.io>

# Where Are We?



# Where Are We?




- ▶ **TensorFlow** is an open source software library for **numerical computation**, particularly well suited and **fine-tuned for large-scale Machine Learning**.
- ▶ Was developed by the **Google Brain team**.



- $$d = a + b$$

# Two Phases of Tensorflow

- ▶ Working with TensorFlow involves **two main phases**.
  1. **Build** a graph
  2. **Execute** it



- ```
import tensorflow as tf

a = tf.constant(5)
b = tf.constant(3)

c = tf.multiply(a, b)
d = tf.add(a, b)
e = tf.add(c, d)
```



- ```
import tensorflow as tf

a = tf.constant(5)
b = tf.constant(3)

c = tf.multiply(a, b)
d = tf.add(a, b)
e = tf.add(c, d)
```

- ```
import tensorflow as tf
```

```

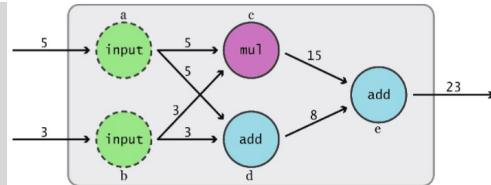
graph LR
    In1(( )) -- 5 --> a((a input))
    In2(( )) -- 3 --> b((b input))
    a -- 5 --> mul((mul))
    a -- 3 --> add1((add))
    b -- 5 --> mul
    b -- 3 --> add1
    mul -- 15 --> add2((e add))
    add1 -- 8 --> add2
    add2 -- 23 --> Out(( ))
    style a fill:#90EE90,stroke-dasharray: 5 5
    style b fill:#90EE90,stroke-dasharray: 5 5
    style mul fill:#DDA0DD
    style add1 fill:#ADD8E6
    style add2 fill:#ADD8E6
  
```

## Phase 2: Execute a Graph

- Now, we are ready to **run the computations**.

```
sess = tf.Session()
print(sess.run(e))
sess.close()


# Alternative way
with tf.Session() as sess:
    print(sess.run(e))
```



- ▶ Now, we are ready to **run the computations**.
- ▶ **Create and run** a **session**, by calling the **run()** method of the **Session** object.

```
sess = tf.Session()
print(sess.run(e))
sess.close()

# Alternative way
with tf.Session() as sess:
    print(sess.run(e))
```




## Phase 2: Execute a Graph

- ▶ Now, we are ready to **run the computations**.
- ▶ **Create and run** a **session**, by calling the **run()** method of the **Session** object.
- ▶ When **sess.run(e)** is called, it starts at the requested output **e**, and then **works backward**, computing nodes that must be executed.

```
sess = tf.Session()
print(sess.run(e))
sess.close()

# Alternative way
with tf.Session() as sess:
    print(sess.run(e))
```




## Phase 2: Execute a Graph

- ▶ Now, we are ready to **run the computations**.
- ▶ **Create and run** a **session**, by calling the **run()** method of the **Session** object.
- ▶ When **sess.run(e)** is called, it starts at the requested output **e**, and then **works backward**, computing nodes that must be executed.
- ▶ **Close the session** at the end of the computation, using the **sess.close()** command.

```
sess = tf.Session()
print(sess.run(e))
sess.close()

# Alternative way
with tf.Session() as sess:
    print(sess.run(e))
```





# The Complete Code

```
import tensorflow as tf

# Building the Graph
a = tf.constant(5)
b = tf.constant(3)

c = tf.multiply(a, b)
d = tf.add(a, b)
e = tf.add(c, d)

# Executing the Graph
with tf.Session() as sess:
    print(sess.run(e))
```



# Visualize the Code

```
import tensorflow as tf

# Building the Graph
a = tf.constant(5)
b = tf.constant(3)

c = tf.multiply(a, b)
d = tf.add(a, b)
e = tf.add(c, d)

writer = tf.summary.FileWriter("./graphs", tf.get_default_graph())

# Executing the Graph
with tf.Session() as sess:
    print(sess.run(e))
```





## Visualize the Code

```
import tensorflow as tf

# Building the Graph
a = tf.constant(5)
b = tf.constant(3)

c = tf.multiply(a, b)
d = tf.add(a, b)
e = tf.add(c, d)

writer = tf.summary.FileWriter("./graphs", tf.get_default_graph())

# Executing the Graph
with tf.Session() as sess:
    print(sess.run(e))

tensorboard --logdir="./graphs" --port 6006
```



## Let's Give Name to Variables

```
import tensorflow as tf

# Building the Graph
a = tf.constant(5, name="a")
b = tf.constant(3, name="b")

c = tf.multiply(a, b, name="c_mul")
d = tf.add(a, b, name="d_add")
e = tf.add(c, d, name="e_add")

writer = tf.summary.FileWriter("./graphs", tf.get_default_graph())

# Executing the Graph
with tf.Session() as sess:
    print(sess.run(e))

tensorboard --logdir="./graphs" --port 6006
```

# Tensor Objects



# What is Tensor?

- ▶ The central **unit of data** in TensorFlow is the **tensor**.
- ▶ An **n-dimensional array** of **primitive values**.



# Tensor Objects

- ▶ The **main object** you manipulate and pass around is the `tf.Tensor`.
- ▶ TensorFlow programs work by **building a graph** of `tf.Tensor` objects, and **running** parts of this graph.
- ▶ Each **Tensor object** is specified by:
  - Rank
  - Shape
  - Datatype

# Tensor Objects - Rank

- ▶ The **number of dimensions**.
  - **rank 0**: **scalar** (number), e.g., 5
  - **rank 1**: **vector**, e.g., [2, 5, 7]
  - **rank 2**: **matrix**, e.g., [[1, 2], [3, 4], [5, 6]]
  - **rank 3**: **3-Tensor**
  - **rank n**: **n-Tensor**
- ▶ The **tf.rank** method determines the **rank** of a **tf.Tensor** object.

```
c = tf.constant([[4], [9], [16], [25]])  
  
r = tf.rank(c) # rank 2
```

## Tensor Objects - Shape

- ▶ The **number of elements** in **each dimension**.
- ▶ The **get\_shape()** returns the **shape of the tensor** as a tuple of integers.

```
c = tf.constant([[[1, 2, 3], [4, 5, 6]],  
                [[1, 1, 1], [2, 2, 2]]])  
  
s = c.get_shape() # (2, 2, 3)
```



## Tensor Objects - Data Types (1/2)

- ▶ We can **explicitly** choose the **data type** a Tensor object.
- ▶ Make sure the data types **match throughout** the graph.
- ▶ We can use the **tf.cast()** method to **change the data type** of a Tensor object.

```
c = tf.constant(4.0, dtype=tf.float64)
x = tf.constant([1, 2, 3], dtype=tf.float32)
y = tf.cast(x, tf.int64)
```



## Tensor Objects - Data Types (2/2)

| Data type     | Python type                | Description                                                                  |
|---------------|----------------------------|------------------------------------------------------------------------------|
| DT_FLOAT      | <code>tf.float32</code>    | 32-bit floating point.                                                       |
| DT_DOUBLE     | <code>tf.float64</code>    | 64-bit floating point.                                                       |
| DT_INT8       | <code>tf.int8</code>       | 8-bit signed integer.                                                        |
| DT_INT16      | <code>tf.int16</code>      | 16-bit signed integer.                                                       |
| DT_INT32      | <code>tf.int32</code>      | 32-bit signed integer.                                                       |
| DT_INT64      | <code>tf.int64</code>      | 64-bit signed integer.                                                       |
| DT_UINT8      | <code>tf.uint8</code>      | 8-bit unsigned integer.                                                      |
| DT_UINT16     | <code>tf.uint16</code>     | 16-bit unsigned integer.                                                     |
| DT_STRING     | <code>tf.string</code>     | Variable-length byte array. Each element of a Tensor is a byte array.        |
| DT_BOOL       | <code>tf.bool</code>       | Boolean.                                                                     |
| DT_COMPLEX64  | <code>tf.complex64</code>  | Complex number made of two 32-bit floating points: real and imaginary parts. |
| DT_COMPLEX128 | <code>tf.complex128</code> | Complex number made of two 64-bit floating points: real and imaginary parts. |
| DT_QINT8      | <code>tf.qint8</code>      | 8-bit signed integer used in quantized ops.                                  |
| DT_QINT32     | <code>tf.qint32</code>     | 32-bit signed integer used in quantized ops.                                 |
| DT_QUINT8     | <code>tf.quint8</code>     | 8-bit unsigned integer used in quantized ops.                                |



## Tensor Objects - Name

- ▶ Each **Tensor object** has an **identifying name**.
- ▶ This name is an **intrinsic string name**, not to be confused with the **name of the variable**.

```
c = tf.constant(4.0, dtype=tf.float64, name="input")
```



## Tensor Objects - Name Scopes

- To deal with **large graphs**, we can use **node grouping** to make it **easier to manage**.



## Tensor Objects - Name Scopes

- ▶ To deal with **large graphs**, we can use **node grouping** to make it **easier to manage**.
- ▶ **Hierarchically** group nodes by their **names**, using `tf.name_scope()` together with the `with` clause.

```
with tf.name_scope("myprefix"):  
    c1 = tf.constant(4.0, dtype=tf.int32, name="input1")  
    c2 = tf.constant(4.0, dtype=tf.float64, name="input2")
```



## Tensor Objects - Name Scopes

- ▶ To deal with **large graphs**, we can use **node grouping** to make it **easier to manage**.
- ▶ **Hierarchically** group nodes by their **names**, using `tf.name_scope()` together with the `with` clause.
- ▶ Below, the name of each operation **within the scope** is prefixed with `myprefix/`, e.g., `myprefix/input1`.

```
with tf.name_scope("myprefix"):  
    c1 = tf.constant(4.0, dtype=tf.int32, name="input1")  
    c2 = tf.constant(4.0, dtype=tf.float64, name="input2")
```



# Main Types of Tensors

- ▶ Constants `tf.constant`
- ▶ Variables `tf.Variable`
- ▶ Placeholders `tf.placeholder`

# Constants

## Constants (1/3)

- The value of **constant** Tensor **cannot be changed** in the future.

```
tf.constant(value, dtype=None, shape=None, name="Const", verify_shape=False)

a = tf.constant([[0, 1], [2, 3]], name="b")
b = tf.constant([[4], [9], [16], [25]], name="c")
```



## Constants (2/3)

- The **initialization** should be with a **value**, not with operation.

| TensorFlow operation                                  | Description                                                                                                                                                      |
|-------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>tf.constant(value)</code>                       | Creates a tensor populated with the value or values specified by the argument <i>value</i>                                                                       |
| <code>tf.fill(shape, value)</code>                    | Creates a tensor of shape <i>shape</i> and fills it with <i>value</i>                                                                                            |
| <code>tf.zeros(shape)</code>                          | Returns a tensor of shape <i>shape</i> with all elements set to 0                                                                                                |
| <code>tf.zeros_like(tensor)</code>                    | Returns a tensor of the same type and shape as <i>tensor</i> with all elements set to 0                                                                          |
| <code>tf.ones(shape)</code>                           | Returns a tensor of shape <i>shape</i> with all elements set to 1                                                                                                |
| <code>tf.ones_like(tensor)</code>                     | Returns a tensor of the same type and shape as <i>tensor</i> with all elements set to 1                                                                          |
| <code>tf.random_normal(shape, mean, stddev)</code>    | Outputs random values from a normal distribution                                                                                                                 |
| <code>tf.truncated_normal(shape, mean, stddev)</code> | Outputs random values from a truncated normal distribution (values whose magnitude is more than two standard deviations from the mean are dropped and re-picked) |
| <code>tf.random_uniform(shape, minval, maxval)</code> | Generates values from a uniform distribution in the range [ <i>minval</i> , <i>maxval</i> )                                                                      |
| <code>tf.random_shuffle(tensor)</code>                | Randomly shuffles a tensor along its first dimension                                                                                                             |



## Constants (3/3)

- What's wrong with constants?



## Constants (3/3)

- ▶ What's wrong with constants?
- ▶ Constants are stored in the graph definition.
- ▶ This makes loading graphs expensive when constants are big.



## Constants (3/3)

- ▶ What's wrong with constants?
- ▶ Constants are stored in the graph definition.
- ▶ This makes loading graphs expensive when constants are big.
- ▶ Only use constants for primitive types.
- ▶ Use variables for more data that requires more memory.

# Variables



## Variables

- ▶ A **variable** represents a Tensor whose **value can be changed** by running ops on it.



# Variables

- ▶ A **variable** represents a Tensor whose **value can be changed** by running ops on it.
- ▶ `tf.Variable` is a **class** with several ops.

# Variables

- ▶ A **variable** represents a Tensor whose **value can be changed** by running ops on it.
- ▶ `tf.Variable` is a **class** with several ops.
- ▶ **Create variables** with `tf.get_variable`.
- ▶ `tf.get_variable` returns an **existing variable** with the given parameters if it is available.

```
# not recommended way to make a variable
tf.Variable(<initial-value>, name)

w = tf.Variable([[0, 1], [2, 3]], name="matrix")

# recommended
tf.get_variable(name, shape=None, dtype=tf.float32, initializer=None,
                regularizer=None, trainable=True, collections=None)

w = tf.get_variable("matrix", initializer=tf.constant([[0, 1], [2, 3]]))
```





# Initialize Variables

- ▶ Variables should be initialized before being used.



# Initialize Variables

- ▶ Variables should be **initialized** before being used.
- ▶ Initialize **all variables** at once.

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())
```



# Initialize Variables

- ▶ Variables should be **initialized** before being used.
- ▶ Initialize **all variables** at once.

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())
```

- ▶ Initialize only a **subset of variables**.

```
with tf.Session() as sess:  
    sess.run(tf.variables_initializer([a, b]))
```



# Initialize Variables

- ▶ **Variables** should be **initialized** before being used.
- ▶ Initialize **all variables** at once.

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())
```

- ▶ Initialize only a **subset of variables**.

```
with tf.Session() as sess:  
    sess.run(tf.variables_initializer([a, b]))
```

- ▶ Initialize a **single variable**.

```
w = tf.Variable(tf.zeros([784,10]))  
  
with tf.Session() as sess:  
    sess.run(w.initializer)
```



## Assign Values to Variables (1/3)

- What does it print?

```
w = tf.get_variable("scalar", initializer=tf.constant(2))  
w.assign(100)  
  
with tf.Session() as sess:  
    sess.run(w.initializer)  
    print(sess.run(w))
```

## Assign Values to Variables (1/3)

- What does it print?

```
w = tf.get_variable("scalar", initializer=tf.constant(2))
w.assign(100)

with tf.Session() as sess:
    sess.run(w.initializer)
    print(sess.run(w))
```

- Prints 2, because `w.assign(100)` creates an `assign` op.
- That op `needs to be executed` in a `session` to take effect.

```
w = tf.get_variable("scalar", initializer=tf.constant(2))
assign_op = w.assign(100)

with tf.Session() as sess:
    sess.run(w.initializer)
    sess.run(assign_op)
    print(sess.run(w))
```

## Assign Values to Variables (2/3)

► What does it print?

```
w = tf.get_variable("scalar", initializer=tf.constant(2))
w_times_two = w.assign(2 * w)

with tf.Session() as sess:
    sess.run(w.initializer)
    print(sess.run(w_times_two))
    print(sess.run(w_times_two))
    print(sess.run(w_times_two))
```

## Assign Values to Variables (3/3)

### ► `assign_add()` and `assign_sub()`

```
w = tf.get_variable("scalar", initializer=tf.constant(2))

with tf.Session() as sess:
    sess.run(w.initializer)

    # increment by 10
    print(sess.run(w.assign_add(10)))

    # decrement by 5
    print(sess.run(w.assign_sub(5)))
```



# Placeholders

# Placeholders

- ▶ **Placeholders** are **built-in structures** for **feeding input values**.
- ▶ **Empty variables** that will be **filled with data later on**.
- ▶ **shape=None** means that tensor of **any shape** will be accepted.
  - E.g., **shape=[None, 10]**: a matrix with 10 columns and any number of rows.

```
tf.placeholder(dtype, shape=None, name=None)  
  
x = tf.placeholder(tf.float32, shape=[None, 10])
```



## Feeding Placeholders (1/2)

► What's **wrong** with this code?

```
a = tf.placeholder(tf.float32, shape=[3])  
  
b = tf.constant([5, 5, 5], tf.float32)  
  
c = a + b  
  
with tf.Session() as sess:  
    print(sess.run(c))
```

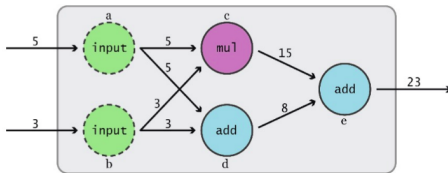
## Feeding Placeholders (2/2)

- Supplement the values to placeholders using a dictionary.


```
a = tf.placeholder(tf.float32, shape=[3])  
b = tf.constant([5, 5, 5], tf.float32)  
  
c = a + b  
  
with tf.Session() as sess:  
    print(sess.run(c, feed_dict={a: [1, 2, 3]}))
```

# Dataflow Graphs

- A computational **graph** is a **series of TensorFlow operations** arranged into a graph.



- ▶ A computational **graph** is a **series of TensorFlow operations** arranged into a graph.
- ▶ The graph is composed of **two types of objects**:
  - **Operations**: the **nodes** of the graph that **consume and produce tensors**.
  - **Tensors**: the **edges** in the graph that **represent the flowing values** through the graph.



# Common Tensorflow Operations

| TensorFlow operator             | Shortcut               | Description                                                                                                            |
|---------------------------------|------------------------|------------------------------------------------------------------------------------------------------------------------|
| <code>tf.add()</code>           | <code>a + b</code>     | Adds a and b, element-wise.                                                                                            |
| <code>tf.multiply()</code>      | <code>a * b</code>     | Multiplies a and b, element-wise.                                                                                      |
| <code>tf.subtract()</code>      | <code>a - b</code>     | Subtracts a from b, element-wise.                                                                                      |
| <code>tf.divide()</code>        | <code>a / b</code>     | Computes Python-style division of a by b.                                                                              |
| <code>tf.pow()</code>           | <code>a ** b</code>    | Returns the result of raising each element in a to its corresponding element b, element-wise.                          |
| <code>tf.mod()</code>           | <code>a % b</code>     | Returns the element-wise modulo.                                                                                       |
| <code>tf.logical_and()</code>   | <code>a &amp; b</code> | Returns the truth table of <code>a &amp; b</code> , element-wise. dtype must be <code>tf.bool</code> .                 |
| <code>tf.greater()</code>       | <code>a &gt; b</code>  | Returns the truth table of <code>a &gt; b</code> , element-wise.                                                       |
| <code>tf.greater_equal()</code> | <code>a &gt;= b</code> | Returns the truth table of <code>a &gt;= b</code> , element-wise.                                                      |
| <code>tf.less_equal()</code>    | <code>a &lt;= b</code> | Returns the truth table of <code>a &lt;= b</code> , element-wise.                                                      |
| <code>tf.less()</code>          | <code>a &lt; b</code>  | Returns the truth table of <code>a &lt; b</code> , element-wise.                                                       |
| <code>tf.negative()</code>      | <code>-a</code>        | Returns the negative value of each element in a.                                                                       |
| <code>tf.logical_not()</code>   | <code>~a</code>        | Returns the logical NOT of each element in a. Only compatible with Tensor objects with dtype of <code>tf.bool</code> . |
| <code>tf.abs()</code>           | <code>abs(a)</code>    | Returns the absolute value of each element in a.                                                                       |
| <code>tf.logical_or()</code>    | <code>a   b</code>     | Returns the truth table of <code>a   b</code> , element-wise. dtype must be <code>tf.bool</code> .                     |





## Managing Multiple Graphs (1/2)

- ▶ When we call `import tensorflow`, a `default graph` is automatically created.



## Managing Multiple Graphs (1/2)

- ▶ When we call `import tensorflow`, a `default graph` is automatically created.
- ▶ We can also create `additional graphs`, by calling `tf.Graph()`.



## Managing Multiple Graphs (1/2)

- ▶ When we call `import tensorflow`, a **default graph** is automatically created.
- ▶ We can also create **additional graphs**, by calling `tf.Graph()`.
- ▶ `tf.get_default_graph()` tells **which graph** is currently set as the **default graph**.

```
import tensorflow as tf

g = tf.Graph()
a = tf.constant(5)

print(a.graph is g)
# Out: False

print(a.graph is tf.get_default_graph())
# Out: True
```

## Managing Multiple Graphs (2/2)

- Use `with` together with `as_default()` to associate your constructed nodes the a right graph.

```
import tensorflow as tf

g1 = tf.get_default_graph()
g2 = tf.Graph()

print(g1 is tf.get_default_graph())
# Out: True

with g2.as_default():
    print(g1 is tf.get_default_graph())
# Out: False
    print(g2 is tf.get_default_graph())
# Out: True
```

# Session



## Session

- ▶ A `Session` object encapsulates the environment.
- ▶ `Operation` objects are `executed`, and `Tensor` objects are `evaluated`.
- ▶ Session will also `allocate memory` to `store` the current values of variables.



# Session

- ▶ A `Session` object encapsulates the environment.
- ▶ `Operation` objects are **executed**, and `Tensor` objects are **evaluated**.
- ▶ `Session` will also **allocate memory** to **store** the current values of variables.

```
sess = tf.Session()
outs = sess.run(e)
print("outs = {}".format(outs))
sess.close()
```



# Session

- ▶ A `Session` object encapsulates the environment.
- ▶ `Operation` objects are *executed*, and `Tensor` objects are *evaluated*.
- ▶ `Session` will also *allocate memory* to *store* the current values of variables.

```
sess = tf.Session()
outs = sess.run(e)
print("outs = {}".format(outs))
sess.close()
```

```
# can be written as follows
with tf.Session() as sess:
    outs = sess.run(e)

print("outs = {}".format(outs))
```



- ▶ A graph can be **parameterized** to accept **external inputs** via **placeholders**.
- ▶ To **feed a placeholder**, the **input data** is passed to the **session.run()**.
- ▶ Each **key** corresponds to a **placeholder variable name**, and the matching **values** are the **data values**.


```
x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = x + y

with tf.Session() as sess:
    print(sess.run(z, feed_dict={x: 3, y: 4.5}))
    print(sess.run(z, feed_dict={x: [1, 3], y: [2, 4]}))
```

# Fetches

- To **fetch** a **list of outputs** of nodes.

```
with tf.Session() as sess:  
    fetches = [a, b, c, d, e]  
    outs = sess.run(fetches)  
  
print("outs = {}".format(outs))
```





## `Session.run()` vs. `Tensor.eval()`

- ▶ Two ways to evaluate part of graph: `Session.run` and `Tensor.eval`.



## Session.run() vs. Tensor.eval()

- ▶ Two ways to **evaluate part of graph**: `Session.run` and `Tensor.eval`.
- ▶ The most important **difference** is that you can use `sess.run()` to fetch the **values of many tensors in the same step**.

```
t = tf.constant(42.0)
u = tf.constant(37.0)
tu = tf.mul(t, u)
ut = tf.mul(u, t)

with sess.as_default():
    tu.eval()  # runs one step
    ut.eval()  # runs one step
    sess.run([tu, ut])  # evaluates both tensors in a single step
```

# Linear Regression in TensorFlow



# Linear Regression

- ▶ We want to find **weights**  $\mathbf{w}$  and a **bias** term  $b$ .
- ▶ Assume our **target value** is a **linear** combination of some input vector  $\mathbf{x}$ :  $\hat{y} = \mathbf{w}^T \mathbf{x} + b$ .



# Linear Regression

- ▶ We want to find **weights**  $\mathbf{w}$  and a **bias** term  $b$ .
- ▶ Assume our **target value** is a **linear** combination of some input vector  $\mathbf{x}$ :  $\hat{y} = \mathbf{w}^T \mathbf{x} + b$ .
- ▶ Let's generate **synthetic data**.

```
import numpy as np
import tensorflow as tf

x_data = np.random.randn(2000, 3)
w_real = [0.3, 0.5, 0.1]
b_real = -0.2
y_data = np.matmul(w_real, x_data.T) + b_real
```

# Linear Regression - Placeholders and Variables

- ▶ Create **placeholders** for our **input and output data**.
- ▶ Create **variables** for our **weights and intercept**.

```
# placeholders
x = tf.placeholder(tf.float32, shape=[None, 3])
y_true = tf.placeholder(tf.float32, shape=None)

# variables
w = tf.get_variable("weights", dtype=tf.float32, initializer=tf.constant([[0., 0., 0.])))
b = tf.get_variable("bias", dtype=tf.float32, initializer=tf.constant(0.))
```





# Linear Regression - Defining a Cost Function

- ▶ We need a good **measure** to evaluate the **model's performance**.
- ▶ Let's define **MSE** (mean squared error).

```
# the cost function  
y_hat = tf.matmul(w, tf.transpose(x)) + b  
  
cost = tf.reduce_mean(tf.square(y_true - y_hat))
```



# Linear Regression - The Gradient Descent Optimizer

- ▶ Next, we need to **minimize the cost function**.
- ▶ Let's use the **gradient descent**.
- ▶ First create an **optimizer** by using the `GradientDescentOptimizer()` function.
- ▶ Then, create a **train operation** by calling the `optimizer.minimize()` to update our variables.

```
# optimizer
learning_rate = 0.5

optimizer = tf.train.GradientDescentOptimizer(learning_rate)
train = optimizer.minimize(cost)
```

# Linear Regression - Execute It

- At the end, we need to initialize the variables and execute the train operation.

```
num_steps = 10
init = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(init)
    for step in range(num_steps):
        sess.run(train, {x: x_data, y_true: y_data})
    print(sess.run([w, b, cost], {x: x_data, y_true: y_data}))
```

# Logistic Regression in TensorFlow

# Logistic Regression

- ▶ We want to find **weights**  $\mathbf{w}$  and a **bias** term  $b$  in a logistic regression model:

$$\hat{y} = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + b)}}$$

- ▶ Let's generate synthetic data.

```
import numpy as np
import tensorflow as tf

x_data = np.random.randn(2000, 3)
w_real = [0.3, 0.5, 0.1]
b_real = -0.2
y_data = np.matmul(w_real, x_data.T) + b_real
```

# Logistic Regression - Placeholders and Variables

- ▶ Create **placeholders** for our **input and output data**.
- ▶ Create **variables** for our **weights and intercept**.

```
# placeholders
x = tf.placeholder(tf.float32, shape=[None, 3])
y_true = tf.placeholder(tf.float32, shape=None)

# variables
w = tf.get_variable("weights", dtype=tf.float32, initializer=tf.constant([[0., 0., 0.])))
b = tf.get_variable("bias", dtype=tf.float32, initializer=tf.constant(0.))
```

# Logistic Regression - Defining a Loss Function

- For the cost function, we use the [cross-entropy](#) model.

```
z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = -y_true * tf.log(y_hat) - (1 - y_true) * tf.log(1 - y_hat)
cost = tf.reduce_mean(cost)
```

# Logistic Regression - Defining a Loss Function

- For the cost function, we use the **cross-entropy** model.

```
z = tf.matmul(w, tf.transpose(x)) + b
y_hat = tf.sigmoid(z)

cost = -y_true * tf.log(y_hat) - (1 - y_true) * tf.log(1 - y_hat)
cost = tf.reduce_mean(cost)
```

- Alternatively, we can use a **designated function** by TensorFlow.

```
cost = tf.nn.sigmoid_cross_entropy_with_logits(labels=y_true, logits=y_hat)
cost = tf.reduce_mean(cost)
```





# Logistic Regression - The Gradient Descent Optimizer

- ▶ Similar to linear regression.

```
learning_rate = 0.5  
  
optimizer = tf.train.GradientDescentOptimizer(learning_rate)  
train = optimizer.minimize(cost)
```

# Logistic Regression - Execute It

- Similar to linear regression.

```
num_steps = 10
init = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(init)
    for step in range(num_steps):
        sess.run(train, {x: x_data, y_true: y_data})
    print(sess.run([w, b, cost], {x: x_data, y_true: y_data}))
```

# Saving and Restoring Models



## Saving Models

- ▶ Save a **model's parameters** in disk.
- ▶ Create a **Saver** node at the **end of the construction phase**.
- ▶ Then, in the **execution phase**, call its **save()** method whenever you want to save the model.

## Saving Models

- ▶ Save a **model's parameters** in disk.
- ▶ Create a **Saver** node at the **end of the construction phase**.
- ▶ Then, in the **execution phase**, call its **save()** method whenever you want to save the model.

```
w = tf.Variable([[0, 0, 0]], dtype=tf.float32, name="weights")
[...]  
init = tf.global_variables_initializer()  
saver = tf.train.Saver()  
  
with tf.Session() as sess:  
    sess.run(init)  
    for step in range(num_steps):  
        if step % 100 == 0: # checkpoint every 100 epochs  
            save_path = saver.save(sess, "/tmp/my_model.ckpt")  
            sess.run(train, {x: x_data, y_true: y_data})  
    best_w = sess.run(w)  
    save_path = saver.save(sess, "/tmp/my_model_final.ckpt")
```



## Restoring Models

- ▶ Create a `Saver` node at the **end of the construction phase**.
- ▶ Then, at the **beginning of the execution phase** call the `restore()` method of the `Saver` node.
  - Instead of initializing the variables using the `init` node.

```
with tf.Session() as sess:  
    saver.restore(sess, "/tmp/my_model_final.ckpt")  
    [...]
```

# TensorBoard



## TensorBoard (1/3)

- ▶ TensorFlow provides a utility called **TensorBoard**.
- ▶ To visualize your model, you need to write the **graph definition** and some **training stats** to a **log directory** that TensorBoard will read from.
- ▶ Use a **different log directory** every time you run your program, or else TensorBoard will merge them.




## TensorBoard (2/3)

- ▶ Add the following code at the **very end of the construction phase**.
- ▶ The first line writes the **cost**.
- ▶ The second line creates a **FileWriter** that writes summaries of the graph.
- ▶ Start the **TensorBoard web server** (port 6006): **tensorboard --logdir .**

```
logdir = "."
mse_summary = tf.summary.scalar("MSE", cost)
file_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())
file_writer.close()
```

Main Graph




Auxiliary Nodes



## TensorBoard (3/3)

```
cost_summary = tf.summary.scalar("Loss", cost)
file_writer = tf.summary.FileWriter('.', tf.get_default_graph())
[...]
for step in range(num_steps):
    sess.run(train, {x: x_data, y_true: y_data})
    summary_str = cost_summary.eval(feed_dict={x: x_data, y_true: y_data})
    file_writer.add_summary(summary_str, step)
```

Loss



# Keras

- ▶ **Keras** is a **high-level** API to build and train deep learning models.
- ▶ To get started, import **tf.keras** to your program.

```
import tensorflow as tf
from tensorflow.keras import layers
```





## Keras Layers (1/2)

- ▶ In Keras, you assemble `layers` `tf.keras.layers` to build `models`.
- ▶ A model is (usually) a `graph of layers`.
- ▶ There are many types of layers, e.g., `Dense`, `Conv2D`, `RNN`, ...



## Keras Layers (2/2)

- ▶ Common constructor **parameters**:

```
layers.Dense(64, activation=tf.sigmoid, kernel_regularizer=tf.keras.regularizers.l1(0.01),  
             bias_initializer=tf.keras.initializers.constant(2.0))
```



## Keras Layers (2/2)

- ▶ Common constructor **parameters**:
  - **activation**: the **activation function** for the layer.

```
layers.Dense(64, activation=tf.sigmoid, kernel_regularizer=tf.keras.regularizers.l1(0.01),  
             bias_initializer=tf.keras.initializers.constant(2.0))
```



## Keras Layers (2/2)

► Common constructor **parameters**:

- **activation**: the **activation function** for the layer.
- **kernel\_initializer** and **bias\_initializer**: the **initialization** schemes of the layer's weights.

```
layers.Dense(64, activation=tf.sigmoid, kernel_regularizer=tf.keras.regularizers.l1(0.01),  
             bias_initializer=tf.keras.initializers.constant(2.0))
```



## Keras Layers (2/2)

► Common constructor **parameters**:

- **activation**: the **activation function** for the layer.
- **kernel\_initializer** and **bias\_initializer**: the **initialization** schemes of the layer's weights.
- **kernel\_regularizer** and **bias\_regularizer**: the **regularization** schemes of the layer's weights, e.g., **L1** or **L2**.

```
layers.Dense(64, activation=tf.sigmoid, kernel_regularizer=tf.keras.regularizers.l1(0.01),  
             bias_initializer=tf.keras.initializers.constant(2.0))
```



# Keras Models

- ▶ There are two ways to build Keras models: sequential and functional.



# Keras Models

- ▶ There are two ways to build Keras models: sequential and functional.
- ▶ The sequential API allows you to create models layer-by-layer.



# Keras Models

- ▶ There are two ways to build Keras models: sequential and functional.
- ▶ The sequential API allows you to create models layer-by-layer.
- ▶ The functional API allows you to create models that have a lot more flexibility.
  - You can define models where layers connect to more than just the previous and next layers.

- ▶ You can use `tf.keras.Sequential` to build a **sequential model**.

```
from tensorflow.keras import layers

model = tf.keras.Sequential()

model.add(layers.Dense(64, activation="relu"))
model.add(layers.Dense(64, activation="relu"))
model.add(layers.Dense(10, activation="softmax"))
```

- ▶ You can use `tf.keras.Model` to build a **functional model**.

```
from tensorflow.keras import layers

inputs = tf.keras.Input(shape=(32,))
x = layers.Dense(64, activation="relu")(inputs)
x = layers.Dense(64, activation="relu")(x)
predictions = layers.Dense(10, activation="softmax")(x)

model = tf.keras.Model(inputs=inputs, outputs=predictions)
```



# Training Keras Models

- ▶ Call the `compile` method to **configure the learning process**.
- ▶ `tf.keras.Model.compile` takes **three important arguments**.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])  
model.fit(trainig_data, training_labels, epochs=10, batch_size=32)
```



# Training Keras Models

- ▶ Call the `compile` method to **configure the learning process**.
- ▶ `tf.keras.Model.compile` takes **three important arguments**.
  - `optimizer`: specifies the **training procedure**.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])  
model.fit(trainig_data, training_labels, epochs=10, batch_size=32)
```





# Training Keras Models

- ▶ Call the `compile` method to **configure the learning process**.
- ▶ `tf.keras.Model.compile` takes **three important arguments**.
  - `optimizer`: specifies the **training procedure**.
  - `loss`: the **cost function** to minimize during optimization, e.g., mean square error (`mse`), `categorical_crossentropy`, and `binary_crossentropy`.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])  
model.fit(trainig_data, training_labels, epochs=10, batch_size=32)
```

# Training Keras Models

- ▶ Call the `compile` method to **configure the learning process**.
- ▶ `tf.keras.Model.compile` takes **three important arguments**.
  - `optimizer`: specifies the **training procedure**.
  - `loss`: the **cost function** to minimize during optimization, e.g., mean square error (`mse`), `categorical_crossentropy`, and `binary_crossentropy`.
  - `metrics`: used to **monitor training**.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])  
model.fit(trainig_data, training_labels, epochs=10, batch_size=32)
```

# Training Keras Models

- ▶ Call the `compile` method to **configure the learning process**.
- ▶ `tf.keras.Model.compile` takes **three important arguments**.
  - `optimizer`: specifies the **training procedure**.
  - `loss`: the **cost function** to minimize during optimization, e.g., mean square error (`mse`), `categorical_crossentropy`, and `binary_crossentropy`.
  - `metrics`: used to **monitor training**.
- ▶ Call the `fit` method to **fit the model the training data**.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])  
model.fit(trainig_data, training_labels, epochs=10, batch_size=32)
```



# Evaluate and Predict

- ▶ `tf.keras.Model.evaluate`: evaluate the cost and metrics for the data provided.
- ▶ `tf.keras.Model.predict`: predict the output of the last layer for the data provided.

```
model.evaluate(test_data, test_labels, batch_size=32)
```

```
model.predict(test_data, batch_size=32)
```

# Linear Regression in Keras

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers

x_data = np.random.randn(2000, 3)
w_real = [0.3, 0.5, 0.1]
b_real = -0.2
y_data = np.matmul(w_real, x_data.T) + b_real

model = tf.keras.Sequential([layers.Dense(1, activation="linear")])

model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001),
              loss="mse", metrics=["mae"])

model.fit(x_data, y_data, epochs=100, batch_size=32)

print(model.get_weights())
```

# Logistic Regression in Keras

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers

x_data = ...
y_data = ...

model = tf.keras.Sequential([layers.Dense(1, activation="sigmoid")])

model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001),
              loss="binary_crossentropy", metrics=["accuracy"])

model.fit(x_data, y_data, epochs=100, batch_size=32)

print(model.get_weights())
```

# Summary



# Summary

- ▶ Dataflow graph
- ▶ Tensors: constants, variables, placeholders
- ▶ Session
- ▶ Save and restore models
- ▶ TensorBoard
- ▶ Keras





## Reference

- ▶ Aurélien Géron, Hands-On Machine Learning (Ch. 9, 12)
- ▶ Some slides were derived from Chip Huyen's slides:  
<http://web.stanford.edu/class/cs20si>

Questions?